

Basic Network & CNN

May 19, 2024

Implementation of `activations.ReLU`:

```
class ReLU(Activation):
    def __init__(self):
        super().__init__()

    def forward(self, Z: np.ndarray) -> np.ndarray:
        """Forward pass for relu activation:
         $f(z) = z$  if  $z \geq 0$ 
        0 otherwise

        Parameters
        -----
        Z    input pre-activations (any shape)

        Returns
        -----
         $f(z)$  as described above applied elementwise to `Z`
        """
        ### YOUR CODE HERE ###
        return np.maximum(0, Z)

    def backward(self, Z: np.ndarray, dY: np.ndarray) -> np.ndarray:
        """Backward pass for relu activation.

        Parameters
        -----
        Z    input to `forward` method
        dY    gradient of loss w.r.t. the output of this layer
             same shape as `Z`

        Returns
        -----
        gradient of loss w.r.t. input of this layer
        """
        ### YOUR CODE HERE ###
        dZ = np.array(Z, copy=True)
        dZ[Z < 0], dZ[Z >= 0] = 0, 1
        return dY * dZ
```

Implementation of layers.FullyConnected:

```
class FullyConnected(Layer):
    """A fully-connected layer multiplies its input by a weight matrix, adds
    a bias, and then applies an activation function.
    """

    def __init__(
        self, n_out: int, activation: str, weight_init="xavier_uniform"
    ) -> None:

        super().__init__()
        self.n_in = None
        self.n_out = n_out
        self.activation = initialize_activation(activation)

        # instantiate the weight initializer
        self.init_weights = initialize_weights(weight_init, activation=activation)

    def _init_parameters(self, X_shape: Tuple[int, int]) -> None:
        """Initialize all layer parameters (weights, biases)."""
        self.n_in = X_shape[1]

        ### BEGIN YOUR CODE ###

        W = self.init_weights((self.n_in, self.n_out)) # Takes the shape of the design matrix
        b = np.zeros((1, self.n_out)) # Initialize bias to zero, 1xn^(l+1)
        Z = []
        X = []

        self.parameters = OrderedDict({"W": W, "b": b}) # DO NOT CHANGE THE KEYS
        self.cache = OrderedDict({"Z": Z, "X": X}) # cache for backprop
        self.gradients = OrderedDict({"W": np.zeros_like(W), "b": np.zeros_like(b)}) # parameters for backprop
        # MUST HAVE THE SAME SHAPES AS THE PARAMETERS

        ### END YOUR CODE ###

    def forward(self, X: np.ndarray) -> np.ndarray:
        """Forward pass: multiply by a weight matrix, add a bias, apply activation.
        Also, store all necessary intermediate results in the `cache` dictionary
        to be able to compute the backward pass.

        Parameters
        -----
        X input matrix of shape (batch_size, input_dim)

        Returns
        -----
        a matrix of shape (batch_size, output_dim)"""
```

```

"""
# initialize layer parameters if they have not been initialized
if self.n_in is None:
    self._init_parameters(X.shape)

### BEGIN YOUR CODE ###
W, b = self.parameters["W"], self.parameters["b"]
# perform an affine transformation and activation
Z = X @ W + b
out = self.activation(Z)
# store information necessary for backprop in `self.cache`
self.cache["Z"], self.cache["X"] = Z, X
### END YOUR CODE ###

return out

def backward(self, dLdY: np.ndarray) -> np.ndarray:
    """Backward pass for fully connected layer.
    Compute the gradients of the loss with respect to:
        1. the weights of this layer (mutate the `gradients` dictionary)
        2. the bias of this layer (mutate the `gradients` dictionary)
        3. the input of this layer (return this)

    Parameters
    -----
    dLdY    gradient of the loss with respect to the output of this layer
            shape (batch_size, output_dim)

    Returns
    -----
    gradient of the loss with respect to the input of this layer
    shape (batch_size, input_dim)
    """
    ### BEGIN YOUR CODE ###

    # unpack the cache
    W, b = self.parameters["W"], self.parameters["b"]
    Z, X = self.cache["Z"], self.cache["X"]
    # compute the gradients of the loss w.r.t. all parameters as well as the
    # input of the layer
    dLdZ = self.activation.backward(Z, dLdY)
    dLdW = X.T @ dLdZ
    dX = dLdZ @ W.T
    dLdb = np.sum(dLdZ, axis=0, keepdims=True)
    # store the gradients in `self.gradients`
    # the gradient for self.parameters["W"] should be stored in
    # self.gradients["W"], etc.
    self.gradients["W"], self.gradients["b"] = dLdW, dLdb

```

```
### END YOUR CODE ###
```

```
return dX
```

Implementation of activations.SoftMax:

```
class SoftMax(Activation):
    def __init__(self):
        super().__init__()

    def forward(self, Z: np.ndarray) -> np.ndarray:
        """Forward pass for softmax activation.
        Hint: The naive implementation might not be numerically stable.

        Parameters
        -----
        Z   input pre-activations (any shape)

        Returns
        -----
        f(z) as described above applied elementwise to `Z`
        """
        ### YOUR CODE HERE ###
        shifted_Z = Z - np.max(Z, axis=-1, keepdims=True)
        exp = np.exp(shifted_Z)
        return exp / np.sum(exp, axis=-1, keepdims=True)

    def backward(self, Z: np.ndarray, dY: np.ndarray) -> np.ndarray:
        """Backward pass for softmax activation.

        Parameters
        -----
        Z   input to `forward` method
        dY  gradient of loss w.r.t. the output of this layer
            same shape as `Z`

        Returns
        -----
        gradient of loss w.r.t. input of this layer
        """
        ### YOUR CODE HERE ###
        softmax = self.forward(Z)
        dZ = dY - np.sum(dY * softmax, axis=-1, keepdims=True)
        return dZ * softmax
```

Implementation of losses.CrossEntropy:

```
class CrossEntropy(Loss):
    """Cross entropy loss function."""
```

```

def __init__(self, name: str) -> None:
    self.name = name

def __call__(self, Y: np.ndarray, Y_hat: np.ndarray) -> float:
    return self.forward(Y, Y_hat)

def forward(self, Y: np.ndarray, Y_hat: np.ndarray) -> float:
    """Computes the loss for predictions `Y_hat` given one-hot encoded labels
    `Y`.

    Parameters
    -----
    Y        one-hot encoded labels of shape (batch_size, num_classes)
    Y_hat    model predictions in range (0, 1) of shape (batch_size, num_classes)

    Returns
    -----
    a single float representing the loss
    """
    ### YOUR CODE HERE ###
    return -(np.sum(Y * np.log(Y_hat))) / Y.shape[0]

def backward(self, Y: np.ndarray, Y_hat: np.ndarray) -> np.ndarray:
    """Backward pass of cross-entropy loss.
    NOTE: This is correct ONLY when the loss function is SoftMax.

    Parameters
    -----
    Y        one-hot encoded labels of shape (batch_size, num_classes)
    Y_hat    model predictions in range (0, 1) of shape (batch_size, num_classes)

    Returns
    -----
    the gradient of the cross-entropy loss with respect to the vector of
    predictions, `Y_hat`
    """
    ### YOUR CODE HERE ###
    return -((Y/Y_hat)/Y.shape[0])

```

Implementation of models.NeuralNetwork.forward:

```

def forward(self, X: np.ndarray) -> np.ndarray:
    """One forward pass through all the layers of the neural network.

    Parameters
    -----
    X    design matrix whose must match the input shape required by the
         first layer

```

Returns

forward pass output, matches the shape of the output of the last layer
"""

YOUR CODE HERE

Iterate through the network's layers.

for layer **in** self.layers:

 X = layer.forward(X)

return X

Implementation of models.NeuralNetwork.backward:

def backward(self, target: np.ndarray, out: np.ndarray) -> float:

"""One backward pass through all the layers of the neural network.

During this phase we calculate the gradients of the loss with respect to each of the parameters of the entire neural network. Most of the heavy lifting is done by the `backward` methods of the layers, so this method should be relatively simple. Also make sure to compute the loss in this method and NOT in `self.forward`.

Note: Both input arrays have the same shape.

Parameters

target the targets we are trying to fit to (e.g., training labels)

out the predictions of the model on training data

Returns

the loss of the model given the training inputs and targets
"""

YOUR CODE HERE

Compute the loss.

Backpropagate through the network's layers.

loss = self.loss.forward(target, out) *# Y, Y_hat*

dLdY = self.loss.backward(target, out)

for layer **in** reversed(self.layers): *# Backpropagate*

 dLdY = layer.backward(dLdY)

return loss

Implementation of models.NeuralNetwork.predict:

def predict(self, X: np.ndarray, Y: np.ndarray) -> Tuple[np.ndarray, float]:

"""Make a forward and backward pass to calculate the predictions and loss of the neural network on the given data.

Parameters

X input features

Y targets (same length as `X`)

Returns

a tuple of the prediction and loss

"""

YOUR CODE HERE

Do a forward pass. Maybe use a function you already wrote?

Get the loss. Remember that the `backward` function returns the loss.

Y_hat = self.forward(X)

loss = self.backward(Y, Y_hat)

return Y_hat, loss

Implementation of layers.Conv2D.forward:

```
def forward(self, X: np.ndarray) -> np.ndarray:
```

*"""Forward pass for convolutional layer. This layer convolves the input
`X` with a filter of weights, adds a bias term, and applies an activation
function to compute the output. This layer also supports padding and
integer strides. Intermediates necessary for the backward pass are stored
in the cache.*

Parameters

X input with shape (batch_size, in_rows, in_cols, in_channels)

Returns

output feature maps with shape (batch_size, out_rows, out_cols, out_channels)

"""

if self.n_in is None:

self._init_parameters(X.shape)

W = self.parameters["W"]

b = self.parameters["b"]

kernel_height, kernel_width, in_channels, out_channels = W.shape

n_examples, in_rows, in_cols, in_channels = X.shape

kernel_shape = (kernel_height, kernel_width)

BEGIN YOUR CODE

pad_h, pad_w = self.pad

out_rows = (in_rows + 2*pad_h - kernel_height) // self.stride + 1 *# formula from disc*

out_cols = (in_cols + 2*pad_w - kernel_width) // self.stride + 1 *# formula from disc*

Padding for four axis

X_pad = np.pad(X, pad_width=((0, 0), (pad_h, pad_h), (pad_w, pad_w), (0, 0)), mode='constant')

implement a convolutional forward pass

Z = np.empty((n_examples, out_rows, out_cols, out_channels), dtype=X.dtype)

for i in range(out_rows):

for j in range(out_cols):

```

        i_start, j_start = i*self.stride, j*self.stride
        i_end, j_end = i_start+kernel_height, j_start+kernel_width
        X_slice = X_pad[:,i_start:i_end,j_start:j_end,:]
        np.einsum('ijcn, bijc->bn', W, X_slice, out=(Z[:,i,j,:]))
# X_pad: (16, 18, 18, 32)    W: (3, 3, 3, 32)    b: (1, 32)    Z: (16, 16, 16, 32)
Z += b
out = self.activation.forward(Z)
# cache any values required for backprop
self.cache['Z'], self.cache['X'] = Z, X
self.cache['X_pad'] = X_pad
self.cache['out_rows'], self.cache['out_cols'] = out_rows, out_cols
### END YOUR CODE ###

return out

```

Implementation of layers.Conv2D.backward:

```

def backward(self, dLdY: np.ndarray) -> np.ndarray:
    """Backward pass for conv layer. Computes the gradients of the output
    with respect to the input feature maps as well as the filter weights and
    biases.

    Parameters
    -----
    dLdY  gradient of loss with respect to output of this layer
          shape (batch_size, out_rows, out_cols, out_channels)

    Returns
    -----
    gradient of the loss with respect to the input of this layer
    shape (batch_size, in_rows, in_cols, in_channels)
    """
    ### BEGIN YOUR CODE ###
    W, b = self.parameters["W"], self.parameters["b"]
    Z, X = self.cache['Z'], self.cache['X']
    pad_h, pad_w = self.pad
    X_pad = self.cache['X_pad']
    out_rows, out_cols = self.cache['out_rows'], self.cache['out_cols']
    # perform a backward pass
    dLdZ = self.activation.backward(Z, dLdY)
    dX = np.zeros_like(X_pad)

    for i in range(out_rows):
        for j in range(out_cols):
            i_start, j_start = i*self.stride, j*self.stride
            i_end, j_end = i_start+W.shape[0], j_start+W.shape[1]
            X_slice = X_pad[:,i_start:i_end,j_start:j_end,:]
            self.gradients['W'] += np.einsum('bn, bijc->ijcn', dLdZ[:,i,j,:], X_slice)
            dX[:,i_start:i_end,j_start:j_end,:] += np.einsum('bn, ijcn->bijc', dLdZ[:,i,j,],

```



```

self.gradients['b'] = np.sum(dLdZ, axis=(0, 1, 2)).reshape(1, -1)
dX = dX[:, pad_h:pad_h+X.shape[1], pad_w:pad_w+X.shape[2],:] # Remove padding
### END YOUR CODE ###

return dX

```

Implementation of layers.Pool2D.forward:

```

def forward(self, X: np.ndarray) -> np.ndarray:
    """Forward pass: use the pooling function to aggregate local information
    in the input. This layer typically reduces the spatial dimensionality of
    the input while keeping the number of feature maps the same.

    As with all other layers, please make sure to cache the appropriate
    information for the backward pass.

    Parameters
    -----
    X    input array of shape (batch_size, in_rows, in_cols, channels)

    Returns
    -----
    pooled array of shape (batch_size, out_rows, out_cols, channels)
    """
    ### BEGIN YOUR CODE ###
    kernel_height, kernel_width = self.kernel_shape
    batch_size, in_rows, in_cols, channels = X.shape
    pad_h, pad_w = self.pad
    out_rows = (in_rows + 2*pad_h - kernel_height) // self.stride + 1
    out_cols = (in_cols + 2*pad_w - kernel_width) // self.stride + 1

    X_pad = np.pad(X, pad_width=((0, 0), (pad_h, pad_h), (pad_w, pad_w), (0, 0)), mode='constant')
    X_pool = np.zeros((batch_size, out_rows, out_cols, channels))
    # implement the forward pass
    for i in range(out_rows):
        for j in range(out_cols):
            i_start, j_start = i*self.stride, j*self.stride
            i_end, j_end = i_start+kernel_height, j_start+kernel_width
            X_pool[:,i,j,:] = self.pool_fn(X_pad[:,i_start:i_end,j_start:j_end,:], axis=(1, 2))
    # cache any values required for backpro
    self.cache['X_pad'], self.cache['X_shape'] = X_pad, X.shape
    ### END YOUR CODE ###

    return X_pool

```

Implementation of layers.Pool2D.backward:

```

def backward(self, dLdY: np.ndarray) -> np.ndarray:
    """Backward pass for pooling layer.

```

Parameters

*dLdY gradient of loss with respect to the output of this layer
shape (batch_size, out_rows, out_cols, channels)*

Returns

*gradient of loss with respect to the input of this layer
shape (batch_size, in_rows, in_cols, channels)*
"""

BEGIN YOUR CODE

ker_h, ker_w = self.kernel_shape

X_pad = self.cache['X_pad']

batch_size, in_rows, in_cols, channels = self.cache['X_shape']

batch_size, out_rows, out_cols, channels = dLdY.shape

pad_h, pad_w = self.pad

dX = np.zeros_like(X_pad)

perform a backward pass

for i in range(out_rows):

 for j in range(out_cols):

 i_start, j_start = i*self.stride, j*self.stride

 i_end, j_end = i_start+ker_h, j_start+ker_w

 if self.mode == 'max':

 X_i = X_pad[:,i_start:i_end,j_start:j_end,:] *# Partition of X_pad*

 X_i_flat = X_i.reshape(batch_size, ker_h*ker_w, channels) *# Flatten spatial dimensions*

 X_i_indices = self.arg_pool_fn(X_i_flat, axis=1) *# Find the argmax of the flattened spatial dimensions*

 batch_idx, channel_idx = np.indices((batch_size, channels))

Create a mask as same shape as X_i_flat

 mask = np.zeros_like(X_i_flat)

 mask[batch_idx, X_i_indices, channel_idx] = 1 *# Set max position to be 1*

 mask = mask.reshape(batch_size, ker_h, ker_w, channels) *# Reshape back to original shape*

 dX[:,i_start:i_end,j_start:j_end,:] += mask * dLdY[:,i:i+1,j:j+1,:]

 else:

 dX[:,i_start:i_end,j_start:j_end,:] += dLdY[:,i:i+1,j:j+1,:] / (ker_h*ker_w)

dX = dX[:, pad_h:pad_h+in_rows, pad_w:pad_w+in_cols, :] *# Remove padding*

END YOUR CODE

return dX

0.0.1 Activation Function Implementations:

Implementation of activations.Linear:

```
class Linear(Activation):
```

```
    def __init__(self):
```

```
        super().__init__()
```

```

def forward(self, Z: np.ndarray) -> np.ndarray:
    """Forward pass for  $f(z) = z$ .

    Parameters
    -----
    Z    input pre-activations (any shape)

    Returns
    -----
     $f(z)$  as described above applied elementwise to `Z`
    """
    return Z

def backward(self, Z: np.ndarray, dY: np.ndarray) -> np.ndarray:
    """Backward pass for  $f(z) = z$ .

    Parameters
    -----
    Z    input to `forward` method
    dY   gradient of loss w.r.t. the output of this layer
         same shape as `Z`

    Returns
    -----
    gradient of loss w.r.t. input of this layer
    """
    return dY

```

Implementation of activations.Sigmoid:

```

class Sigmoid(Activation):
    def __init__(self):
        super().__init__()

    def forward(self, Z: np.ndarray) -> np.ndarray:
        """Forward pass for sigmoid function:
         $f(z) = 1 / (1 + \exp(-z))$ 

        Parameters
        -----
        Z    input pre-activations (any shape)

        Returns
        -----
         $f(z)$  as described above applied elementwise to `Z`
        """
        ### YOUR CODE HERE ###
        return 1 / (1 + np.exp(-Z))

```

```

def backward(self, Z: np.ndarray, dY: np.ndarray) -> np.ndarray:
    """Backward pass for sigmoid.

    Parameters
    -----
    Z    input to `forward` method
    dY    gradient of loss w.r.t. the output of this layer
          same shape as `Z`

    Returns
    -----
    gradient of loss w.r.t. input of this layer
    """
    ### YOUR CODE HERE ###
    sigmoid = self.forward(Z)
    return dY * (sigmoid * (1-sigmoid))

```

Implementation of activations.ReLU:

```

class ReLU(Activation):
    def __init__(self):
        super().__init__()

    def forward(self, Z: np.ndarray) -> np.ndarray:
        """Forward pass for relu activation:
         $f(z) = z$  if  $z \geq 0$ 
        0 otherwise

        Parameters
        -----
        Z    input pre-activations (any shape)

        Returns
        -----
         $f(z)$  as described above applied elementwise to `Z`
        """
        ### YOUR CODE HERE ###
        return np.maximum(0, Z)

    def backward(self, Z: np.ndarray, dY: np.ndarray) -> np.ndarray:
        """Backward pass for relu activation.

        Parameters
        -----
        Z    input to `forward` method
        dY    gradient of loss w.r.t. the output of this layer
              same shape as `Z`

```

Returns

gradient of loss w.r.t. input of this layer
"""

YOUR CODE HERE

dZ = np.array(Z, copy=True)

dZ[Z<0], dZ[Z>=0] = 0, 1

return dY * dZ